

ASJAL AMIR

DATA ENGINEER

+92-324-2013133 • asjal.amer@gmail.com • github.com/mohdasjal1 • linkedin.com/in/asjal-amir • [Portfolio](#)

PROFESSIONAL SUMMARY

Data Engineer with 2+ years of experience building production-grade ETL pipelines, data modeling, and warehousing solutions using AWS, Azure, Apache Spark, and Snowflake. Skilled in Python and SQL across relational and NoSQL databases, automating ingestion of 10M+ daily records at 99.9% uptime. Seeking to apply ETL/ELT and data modeling expertise to deliver reliable, secure data infrastructure for large-scale digital platforms.

WORK EXPERIENCE

Cloud Data Engineer | *Morpheralabs*

Feb 2025 – Present

- Designed and independently launched ETL pipelines in production using Apache Airflow and Apache Spark, enabling data preprocessing and transformation for 10M+ daily records with 35% lower latency.
- Built data models and data warehousing solutions in Snowflake, improving query speeds by 40% for 5+ cross-functional teams to support analytics and insights generation.
- Developed real-time data workflows on AWS (S3, Lambda, Kafka) to enable smarter decision-making, achieving 99.9% pipeline uptime.

Data Engineer Intern | *VentureDive*

Oct 2024 – Jan 2025

- Configured and optimized AWS Glue ETL components for big data pipelines, automating integration of 5M+ structured and semi-structured records.
- Applied SQL query optimization (distribution/sort keys) on Redshift data warehouse, reducing load times by 45%.
- Collaborated with a cross-functional engineering team via Git version control and Agile workflows to deploy fixes, cutting errors by 25%.

SKILLS

Languages: Python (NumPy, Pandas, Matplotlib, PySpark), SQL, JavaScript, Java

Databases: MySQL, PostgreSQL, MongoDB, Microsoft SQL Server (relational & NoSQL)

Cloud Platforms: AWS (S3, Lambda, Redshift, EC2, Glue), Azure, Snowflake

Data Processing Frameworks: Apache Spark, AWS Glue, Apache Airflow, ETL/ELT tools

Version Control & DevOps: Git, Docker, Kubernetes

Data Formats: JSON, Parquet, CSV

PROJECTS

AWS–NiFi–Snowflake SCD Type 2 Pipeline | Amazon S3, Snowpipe, Snowflake Streams & Tasks, Docker, Python [GitHub Link](#)

- Engineered SCD Type 2 pipeline on Snowflake processing 10K+ customer records with full historical tracking.
- Dockerized Apache NiFi on Amazon EC2 to auto-ingest CSV files from Amazon S3 into Snowflake via Snowpipe.
- Implemented Snowflake Streams and Tasks to execute automated SCD2 MERGE, reducing processing overhead by 40%.

Azure Databricks Medallion Pipeline | ADLS Gen2, Delta Lake, Azure Data Factory, Unity Catalog, PySpark [GitHub Link](#)

- Architected Medallion pipeline on Azure Databricks ingesting 5 years of PSX KSE-100 stock data into ADLS Gen2.
- Engineered Silver-layer PySpark transforms with UPSERT merges and anomaly detection for price swings >15%.
- Designed 4 Gold-layer KPI tables (52-week highs, rolling volatility) orchestrated via Azure Data Factory.

AWS Data Lake Pipeline | S3, AWS Glue, Amazon Athena, Redshift, Python [GitHub Link](#)

- Engineered end-to-end lake-to-warehouse pipeline ingesting 12K+ COVID records from S3 into Redshift.
- Automated schema discovery using AWS Glue Crawler and queried raw data via Amazon Athena pre-load.
- Designed dimensional data model and loaded transformed data into Redshift for analytics reporting.

EDUCATION

Bachelor in Computer Science | University of Karachi

2023 – 2026

CERTIFICATIONS

Associate Data Engineer | Datacamp

2025

Cloud Data Engineer | SMIT (Saylani Mass IT Training)

2025

Data Engineering with dbt | LinkedIn Learning

2025

Hands-on essentials Data Engineering | Snowflake

2025